

NovaTech Solutions Pvt. Ltd.

Incident Response Process

Document ID: OPS-IR-001

Version: 2.0

Prepared By: IT Operations & Platform Engineering

Effective Date: January 2026

Classification: Internal Use Only

Document Control

Field	Value
Document	Incident Response Process
Department	IT Operations
Version	2.0
Owner	Platform Engineering
Review Cycle	Every 6 Months
Last Updated	January 2026

Table of Contents

1. Introduction
 2. Purpose
 3. Scope
 4. Incident Definition
 5. Incident Severity Levels
 6. Roles and Responsibilities
 7. Incident Lifecycle Overview
 8. Revision History
-

1. Introduction

The Incident Response Process defines the standard procedures for identifying, assessing, managing, resolving, and reviewing operational incidents affecting the NovaDesk platform and supporting infrastructure.

A structured response minimizes service disruption, protects customer data, and ensures timely communication throughout the incident lifecycle.

This document applies to:

- Production environments
 - Staging environments (high-impact incidents)
 - Platform infrastructure
 - Enterprise applications
 - Databases
 - Network services
 - Cloud infrastructure
-

2. Purpose

The objectives of this process are to:

- Restore services as quickly as possible.
 - Minimize business impact.
 - Protect customer and company data.
 - Coordinate technical response teams.
 - Standardize incident handling.
 - Ensure clear stakeholder communication.
 - Capture lessons learned for continuous improvement.
-

3. Scope

This document applies to operational incidents involving:

- Application failures
- Infrastructure outages
- Database failures
- Network disruptions
- Authentication issues

- Security incidents
- Performance degradation
- Third-party service failures

Routine service requests and planned maintenance activities are managed through separate operational procedures.

4. Incident Definition

An incident is any unplanned event that interrupts or degrades the normal operation of a service.

Examples include:

- Production application unavailable
- Database outage
- API latency exceeding service targets
- Authentication failures
- Failed production deployment
- Storage system failure
- DNS resolution issues
- Load balancer malfunction
- Security breach
- Cloud infrastructure outage

Incidents are prioritized according to business impact and urgency.

Incident Categories

Standard categories include:

Category	Examples
Application	API failures, UI issues
Infrastructure	Server or Kubernetes failures
Database	Connectivity, corruption, replication issues
Network	Routing, DNS, firewall problems
Security	Unauthorized access, malware, credential compromise

Category	Examples
----------	----------

Third-Party	External API or SaaS service outages
-------------	--------------------------------------

Categorization helps assign incidents to the appropriate response team.

5. Incident Severity Levels

Every incident is assigned a severity level based on business impact.

Priority 1 (P1) – Critical

Characteristics:

- Complete production outage
- Significant customer impact
- Business operations unavailable
- No practical workaround

Examples:

- Production unavailable
- Database failure
- Authentication service outage
- Major security compromise

Target Response Time

15 minutes

Target Resolution Time

4 hours

Priority 2 (P2) – High

Characteristics:

- Major functionality unavailable
- Large number of users affected
- Temporary workaround available

Examples:

- Reporting unavailable

- Major API degradation
- Regional service interruption

Target Response Time

30 minutes

Target Resolution Time

8 hours

Priority 3 (P3) – Medium

Characteristics:

- Limited business impact
- Small group of users affected
- Workaround available

Examples:

- Minor feature malfunction
- Performance degradation
- Individual integration issue

Target Response Time

2 hours

Target Resolution Time

2 business days

Priority 4 (P4) – Low

Characteristics:

- Cosmetic issues
- Minor defects
- Enhancement requests
- Documentation corrections

Examples:

- UI formatting issue
- Non-critical bug

- Documentation error

Target Response Time

1 business day

Target Resolution Time

5 business days

6. Roles and Responsibilities

Effective incident management requires clearly defined responsibilities.

Incident Manager

The Incident Manager is responsible for:

- Declaring incidents.
 - Coordinating response activities.
 - Managing communications.
 - Tracking progress.
 - Confirming incident resolution.
 - Leading post-incident review meetings.
-

Platform Engineering Team

Responsibilities include:

- Infrastructure investigation.
 - Kubernetes troubleshooting.
 - Deployment rollback.
 - System recovery.
 - Performance analysis.
-

Database Administration Team

Responsible for:

- Database diagnostics.
- Backup restoration.

- Replication monitoring.
 - Database recovery.
 - Query performance analysis.
-

Information Security Team

Responsibilities include:

- Security incident investigation.
 - Threat containment.
 - Credential management.
 - Digital evidence preservation.
 - Regulatory notification guidance.
-

Application Support Team

Responsible for:

- Application diagnostics.
 - Log analysis.
 - Configuration review.
 - Functional validation.
 - User communication support.
-

Communications Coordinator

Responsible for:

- Internal updates.
- Executive notifications.
- Customer communications.
- Status page updates.
- Incident summaries.

Only authorized personnel may issue external communications.

7. Incident Lifecycle Overview

Every incident follows a standardized lifecycle.

Detection

|



Reporting

|



Assessment

|



Classification

|



Assignment

|



Investigation

|



Containment

|



Recovery

|



Validation

|



Closure

|



Post-Incident Review

Each stage must be documented within the Incident Management System.

Incident Status Values

Standard incident states include:

Status	Description
New	Incident reported
Acknowledged	Response team assigned
Investigating	Root cause analysis in progress
Contained	Immediate impact controlled
Recovering	Service restoration underway
Monitoring	Service validated after recovery
Resolved	Technical resolution completed
Closed	Incident documentation completed

Status updates should be recorded promptly throughout the incident lifecycle.

Incident Records

Each incident record should contain:

- Incident ID
- Date and time reported
- Reporter
- Severity level
- Category
- Affected services
- Impact description
- Assigned response team
- Timeline of events
- Resolution summary

- Closure date

Incident records support operational reporting, audits, and continuous improvement initiatives.

Revision History

Version	Date	Description
1.0	January 2024	Initial Incident Response Process
1.5	July 2025	Added Severity Matrix and Incident Lifecycle
2.0	January 2026	Updated Roles, Response Targets, and Incident Classification

8. Incident Detection

Incidents may be detected through automated monitoring, user reports, or operational teams. Early detection minimizes business impact and accelerates service restoration.

Detection Sources

Incidents may originate from:

- Infrastructure monitoring alerts
- Application monitoring
- Database monitoring
- Security monitoring
- Customer support tickets
- Employee reports
- Third-party service notifications
- Scheduled health checks

Multiple detection sources improve operational visibility.

Automated Monitoring

The monitoring platform continuously evaluates:

- Application availability
- API response times
- CPU utilization

- Memory utilization
- Disk utilization
- Database connectivity
- Queue length
- Error rates
- Certificate expiration

Critical alerts automatically notify the on-call engineering team.

9. Incident Reporting

All incidents should be reported immediately through the approved Incident Management System.

Required Information

Incident reports should include:

- Reporter name
- Date and time detected
- Affected service
- Business impact
- Initial symptoms
- Error messages (if available)
- Severity estimate
- Supporting screenshots or logs (where applicable)

Incomplete reports may delay investigation.

Reporting Channels

Approved reporting channels include:

- IT Service Desk
- Incident Management Portal
- On-call support hotline
- Monitoring platform alerts

Critical production incidents should be reported by telephone in addition to electronic reporting.

10. Incident Assessment

Following initial reporting, the Incident Manager performs an assessment.

Assessment Objectives

Determine:

- Business impact
- Number of affected users
- Service availability
- Potential security implications
- Estimated urgency
- Initial severity level

Severity may be adjusted as additional information becomes available.

Initial Triage

During triage, responders should:

- Verify the reported issue.
- Confirm affected services.
- Identify recent changes.
- Review monitoring dashboards.
- Check application logs.
- Determine whether similar incidents already exist.

Duplicate incidents should be linked to the primary incident record.

11. Investigation

Once the incident has been classified, technical teams begin investigation.

Investigation Activities

Typical activities include:

- Log analysis
- Infrastructure review
- Database diagnostics
- Network diagnostics
- Deployment verification
- Configuration review
- Performance analysis
- Dependency validation

Investigation should be documented throughout the incident.

Evidence Collection

Investigators should collect:

- System logs
- Monitoring graphs
- Error messages
- Configuration snapshots
- Deployment records
- Audit logs

Evidence should be preserved for future analysis.

12. Containment

Containment limits business impact while permanent resolution is prepared.

Containment Techniques

Examples include:

- Isolating affected systems
- Blocking malicious traffic
- Disabling faulty services
- Rolling back recent deployments
- Redirecting traffic

- Activating failover systems

Containment actions should prioritize service stability.

Temporary Workarounds

Where immediate resolution is not possible, temporary workarounds may be implemented.

Examples:

- Read-only mode
- Feature disablement
- Traffic throttling
- Alternate service endpoints

Workarounds should be documented and communicated to stakeholders.

13. Eradication

After containment, the root technical issue should be removed.

Common Eradication Activities

- Remove faulty configuration
- Replace defective application build
- Apply security patches
- Remove malicious software
- Correct database inconsistencies
- Repair infrastructure configuration

Eradication activities should be validated before recovery begins.

14. Recovery

Recovery restores affected services to normal operation.

Recovery Steps

Standard recovery process:

1. Restore infrastructure.

2. Restart affected services.
3. Restore database functionality.
4. Validate integrations.
5. Verify customer access.
6. Monitor service stability.

Recovery should follow approved operational procedures.

Validation

Before closing recovery:

Verify:

- Application availability
- Authentication
- Database connectivity
- API functionality
- Background processing
- Monitoring alerts
- Customer workflows

Successful validation is required before declaring service restored.

15. Communication Procedures

Clear communication reduces uncertainty during incidents.

Internal Communication

Stakeholders should receive updates including:

- Incident summary
- Business impact
- Current status
- Planned actions
- Estimated recovery timeline
- Next update schedule

Updates should be factual and avoid speculation.

Customer Communication

Customer notifications may include:

- Service disruption notice
- Current impact
- Available workarounds
- Progress updates
- Resolution confirmation

Only authorized Communications Coordinators may publish customer-facing updates.

Communication Frequency

Severity Update Frequency

P1	Every 30 Minutes
P2	Every 60 Minutes
P3	Every 4 Hours
P4	As Needed

Additional updates should be provided whenever significant progress occurs.

16. Escalation Matrix

Escalation ensures appropriate expertise is engaged promptly.

Technical Escalation

Severity	Initial Team	Escalation Target
P1	On-call Engineer	Platform Engineering Lead
P2	Application Support	Technical Lead
P3	Support Engineer	Department Manager
P4	Service Desk	Product Owner (if required)

Management Escalation

For major incidents:

1. Incident Manager
2. Platform Engineering Manager
3. Director of Engineering
4. Executive Leadership (if business-critical)

Management escalation is based on business impact rather than technical complexity alone.

External Escalation

External vendors may be engaged for:

- Cloud infrastructure issues
- Third-party SaaS failures
- Internet connectivity problems
- Security services
- Hardware replacement

Vendor engagement should follow contractual support procedures.

17. Incident Documentation

Throughout the incident, responders should record:

- Timeline of events
- Investigation findings
- Decisions made
- Containment actions
- Recovery activities
- Communications issued
- Final resolution

Accurate documentation supports post-incident review and compliance.

18. Operational Best Practices

Incident responders should:

- Acknowledge incidents promptly.
- Assign clear ownership.
- Document all actions.
- Communicate regularly.
- Prioritize customer impact.
- Verify recovery before closure.
- Preserve investigation evidence.
- Escalate without unnecessary delay.
- Avoid making unverified assumptions.
- Maintain a complete audit trail.

Consistent execution improves response effectiveness and organizational learning.

19. Post-Incident Review

Every Priority 1 (P1) and Priority 2 (P2) incident requires a formal Post-Incident Review (PIR).

Priority 3 (P3) and Priority 4 (P4) incidents may also require review when recurring patterns or significant operational risks are identified.

The objective is organizational learning rather than assigning blame.

Review Timeline

Severity Review Target

P1	Within 3 Business Days
P2	Within 5 Business Days
P3	As Required
P4	As Required

Participants

The review should include:

- Incident Manager
- Platform Engineering
- Application Support

- Database Administration
- Information Security (where applicable)
- Product Owner
- Technical Lead

Additional stakeholders may participate depending on the nature of the incident.

20. Root Cause Analysis (RCA)

Root Cause Analysis identifies the underlying reason for the incident to reduce the likelihood of recurrence.

RCA Objectives

The investigation should determine:

- What happened?
- When did it happen?
- Why did it happen?
- How was it detected?
- Why was it not prevented?
- Which controls succeeded?
- Which controls failed?
- What corrective actions are required?

The analysis should focus on systems, processes, and controls rather than individuals.

Investigation Techniques

Common approaches include:

- Five Whys
- Timeline Analysis
- Fault Tree Analysis
- Change Review
- Log Analysis
- Dependency Mapping

Supporting evidence should be retained with the incident record.

RCA Deliverables

The completed analysis should include:

- Incident summary
- Business impact
- Timeline of events
- Technical findings
- Root cause
- Contributing factors
- Corrective actions
- Preventive actions
- Responsible owners
- Target completion dates

21. Corrective and Preventive Actions (CAPA)

Following the RCA, corrective and preventive actions are assigned.

Corrective Actions

Examples include:

- Fix application defects.
- Replace faulty infrastructure.
- Correct configuration errors.
- Restore missing monitoring.
- Improve documentation.

Corrective actions address the immediate cause.

Preventive Actions

Examples include:

- Introduce automated testing.
- Add monitoring alerts.

- Improve deployment validation.
- Increase system redundancy.
- Update operational procedures.
- Provide additional staff training.

Preventive actions reduce future operational risk.

Action Tracking

Each action should include:

- Action ID
- Description
- Owner
- Priority
- Due Date
- Current Status
- Completion Evidence

Open actions are reviewed during operational governance meetings.

22. Incident Metrics and KPIs

Operational metrics are used to evaluate the effectiveness of incident management.

Core Metrics

The following metrics are monitored:

Metric	Description
Mean Time to Detect (MTTD)	Time from occurrence to detection
Mean Time to Acknowledge (MTTA)	Time until response begins
Mean Time to Resolve (MTTR)	Time until service restoration
Incident Volume	Number of incidents reported
Repeat Incidents	Recurring operational issues
SLA Compliance	Achievement of response and resolution targets

Metric	Description
Customer Impact	Number of affected users or services

Metrics are reviewed monthly by IT Operations.

Reporting

Monthly operational reports summarize:

- Incident trends
- Severity distribution
- Response performance
- Resolution performance
- Major incidents
- Outstanding corrective actions
- Recommendations for improvement

Reports support management decision-making and resource planning.

23. Documentation Requirements

Accurate documentation is essential for compliance, auditing, and organizational learning.

Required Records

Each incident should contain:

- Incident identifier
- Date and time
- Reporter
- Assigned responders
- Severity
- Business impact
- Investigation notes
- Timeline
- Communications issued
- Resolution summary

- Root Cause Analysis (if required)
- Corrective actions
- Closure approval

Documentation should be completed before final closure.

Record Retention

Incident records are retained according to organizational record retention policies.

Priority 1 and Priority 2 incident records should remain available for long-term operational analysis and audit purposes.

24. Continuous Improvement

Incident management is reviewed regularly to improve operational maturity.

Improvement Activities

Regular activities include:

- Reviewing recurring incidents.
- Updating operational runbooks.
- Improving monitoring coverage.
- Automating manual tasks.
- Enhancing alert quality.
- Reviewing escalation effectiveness.
- Conducting simulation exercises.

Continuous improvement reduces operational risk over time.

Knowledge Sharing

Lessons learned should be shared through:

- Engineering meetings
- Internal Knowledge Base articles
- Technical documentation updates
- Operational training sessions

Sharing operational knowledge improves future incident response.

25. Compliance and Governance

Incident management activities should align with organizational governance requirements.

Governance Principles

The incident management process should ensure:

- Defined ownership
- Documented procedures
- Management oversight
- Audit readiness
- Evidence preservation
- Continuous review

Governance responsibilities are shared between IT Operations, Platform Engineering, Information Security, and executive leadership.

Operational Audits

Periodic audits review:

- Incident documentation quality
- SLA compliance
- Corrective action completion
- Process adherence
- Communication effectiveness

Audit findings are tracked until resolved.

26. Operational Best Practices

Incident response teams should:

- Detect incidents as early as possible.
- Prioritize customer impact.
- Maintain accurate timelines.
- Escalate promptly when required.
- Communicate consistently.

- Preserve investigation evidence.
- Validate recovery thoroughly.
- Document every major decision.
- Complete post-incident reviews on schedule.
- Implement preventive improvements.

Consistent operational discipline improves service reliability and organizational resilience.

27. Frequently Asked Questions

Q1. Which incidents require a formal Post-Incident Review?

All **Priority 1** and **Priority 2** incidents.

Q2. What is the purpose of Root Cause Analysis?

To identify the underlying cause of the incident and define corrective and preventive actions that reduce the likelihood of recurrence.

Q3. What does MTTR represent?

Mean Time to Resolve — the elapsed time required to restore normal service.

Q4. What information should be included in a completed incident record?

Incident details, investigation findings, communications, resolution summary, and corrective actions.

Q5. Should incident reviews assign blame?

No. Reviews focus on systems, processes, and operational improvements rather than individuals.

Q6. Who owns corrective actions?

Each action is assigned to a named owner with a defined priority and target completion date.

Q7. How are recurring incidents managed?

Recurring incidents are analyzed during operational reviews, and preventive improvements are planned to address underlying causes.

Q8. Why are simulation exercises performed?

Simulation exercises validate response procedures, improve team readiness, and identify opportunities to strengthen the incident management process.

Q9. How often are incident management metrics reviewed?

Core operational metrics are reviewed monthly.

Q10. What is the final step before an incident is closed?

Confirm that documentation is complete, corrective actions have been identified (where required), and the closure has been approved according to the incident management process.

Revision History

Version	Date	Description
1.0	January 2024	Initial Incident Response Process
1.5	July 2025	Added Post-Incident Review and Root Cause Analysis
2.0	January 2026	Added Operational Metrics, Governance, Continuous Improvement, and Documentation Standards

Related Documents

- Infrastructure Guide (TECH-INFRA-001)
 - Database Backup Manual (TECH-DB-001)
 - Developer Guide (TECH-DEV-001)
 - Information Security Policy (COMP-SEC-003)
 - Business Continuity Plan (OPS-BCP-002)
 - Disaster Recovery Runbook (OPS-DR-001)
-

Footer (on every page)

NovaTech Incident	Solutions Response	Pvt.	Ltd. Process
----------------------	-----------------------	------	-----------------

